

## MLP on MNIST Dataset

In [1]:

```
import keras
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
import matplotlib.pyplot as plt
import numpy as np
import time
```

Using TensorFlow backend.

## [1] Dataset Loading and pre-processing

In [2]:

```
# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

In [3]:

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)  
 Number of training examples : 10000 and each image is of shape (28, 28)

In [4]:

```
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

In [5]:

```
# after converting the input images from 3d to 2d vectors

print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)"%(X_train.shape[1]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)  
 Number of training examples : 10000 and each image is of shape (784)

In [6]:

```
# An example data point
print(X_train[0])
```

[illegible]

In [7]:

[illegible]

0.	0.	0.	0.	0.	0.
0.	0.	0.01176471	0.07058824	0.07058824	0.07058824
0.49411765	0.53333333	0.68627451	0.10196078	0.65098039	1.
0.96862745	0.49803922	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.11764706	0.14117647	0.36862745	0.60392157
0.66666667	0.99215686	0.99215686	0.99215686	0.99215686	0.99215686
0.88235294	0.6745098	0.99215686	0.94901961	0.76470588	0.25098039
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.19215686
0.93333333	0.99215686	0.99215686	0.99215686	0.99215686	0.99215686
0.99215686	0.99215686	0.99215686	0.98431373	0.36470588	0.32156863
0.32156863	0.21960784	0.15294118	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.07058824	0.85882353	0.99215686
0.99215686	0.99215686	0.99215686	0.99215686	0.77647059	0.71372549
0.96862745	0.94509804	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.31372549	0.61176471	0.41960784	0.99215686
0.99215686	0.80392157	0.04313725	0.	0.16862745	0.60392157
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.05490196	0.00392157	0.60392157	0.99215686	0.35294118
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.54509804	0.99215686	0.74509804	0.00784314	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.04313725
0.74509804	0.99215686	0.2745098	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.1372549	0.94509804
0.88235294	0.62745098	0.42352941	0.00392157	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.31764706	0.94117647	0.99215686
0.99215686	0.46666667	0.09803922	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.17647059	0.72941176	0.99215686	0.99215686
0.58823529	0.10588235	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.0627451	0.36470588	0.98823529	0.99215686	0.73333333
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.97647059	0.99215686	0.97647059	0.25098039	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.18039216	0.50980392	0.71764706	0.99215686
0.99215686	0.81176471	0.00784314	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.15294118	0.58039216
0.89803922	0.99215686	0.99215686	0.99215686	0.98039216	0.71372549
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.09411765	0.44705882	0.86666667	0.99215686	0.99215686	0.99215686
0.99215686	0.78823529	0.30588235	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.09019608	0.25882353	0.83529412	0.99215686
0.99215686	0.99215686	0.99215686	0.77647059	0.31764706	0.00784314

[illegible]

In [8]:

```
# here we are having a class number for each image
print("Class label of first image :", y_train[0])

# lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])
```

```
Class label of first image : 5
After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
```

## [2] MLP Models

In [9]:

```
from keras.models import Sequential
from keras.layers import Dense, Activation
from keras.layers import Dropout
from keras.layers.normalization import BatchNormalization
```

In [10]:

```
# some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 30
```

In [11]:

```
def plt_epoch_vs_loss(x, vy, ty):
    fig = plt.figure(figsize=(9, 7))
```

```
fig = plt.figure(figsize=(5, 7))
sns.set_style("whitegrid", {'axes.grid' : True})
plt.plot(x, vy, 'b', label="Validation Loss")
plt.plot(x, ty, 'r', label="Train Loss")
plt.legend()
plt.xlabel("epoch")
plt.ylabel("Categorical Crossentropy Loss")
plt.title("Loss")
plt.show()
```

## [2.1] Model 1

Input(784) - ReLu(BatchNormalization(256)) - Dropout(0.5) - ReLu(BatchNormalization(128)) - Dropout(0.25) - Softmax(Output(10)) - Adam Optimizer

In [12]:

```
# for relu layers we are using 'He-Normal weight Initialization'
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with
 $\sigma = \sqrt{2/(fan\_in)}$ .
# h1 =>  $\sigma = \sqrt{2/(256)} = 0.088 \Rightarrow N(0, \sigma) = N(0, 0.088)$ 
# h2 =>  $\sigma = \sqrt{2/(128)} = 0.125 \Rightarrow N(0, \sigma) = N(0, 0.125)$ 

model1 = Sequential()
model1.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(
    mean=0.0, stddev=0.088, seed=None)))
model1.add(BatchNormalization())
model1.add(Dropout(0.5))

model1.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, se
    ed=None)) )
model1.add(BatchNormalization())
model1.add(Dropout(0.25))

model1.add(Dense(output_dim, activation='softmax'))

model1.summary()
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 256)	200960
batch_normalization_1 (Batch Normalization)	(None, 256)	1024
dropout_1 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 128)	32896
batch_normalization_2 (Batch Normalization)	(None, 128)	512
dropout_2 (Dropout)	(None, 128)	0
dense_3 (Dense)	(None, 10)	1290
Total params: 236,682		
Trainable params: 235,914		
Non-trainable params: 768		

In [13]:

```
model1.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model1.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

WARNING:tensorflow:From C:\Users\sanjeev\Anaconda3\lib\site-packages\keras\backend\tensorflow\_backend.py:422: The name tf.global\_variables is deprecated. Please use tf.compat.v1.global\_variables instead.

Train on 60000 samples, validate on 10000 samples

Epoch 1/30

60000/60000 [=====] - 6s 97us/step - loss: 0.4985 - accuracy: 0.8469 - val\_loss: 0.1711 - val\_accuracy: 0.9486

Epoch 2/30

60000/60000 [=====] - 4s 73us/step - loss: 0.2377 - accuracy: 0.9275 - val\_loss: 0.1235 - val\_accuracy: 0.9611

Epoch 3/30

60000/60000 [=====] - 4s 70us/step - loss: 0.1889 - accuracy: 0.9424 - val\_loss: 0.1017 - val\_accuracy: 0.9688

Epoch 4/30

60000/60000 [=====] - 4s 64us/step - loss: 0.1571 - accuracy: 0.9520 - val\_loss: 0.0929 - val\_accuracy: 0.9706

Epoch 5/30

60000/60000 [=====] - 4s 66us/step - loss: 0.1395 - accuracy: 0.9567 - val\_loss: 0.0855 - val\_accuracy: 0.9735

Epoch 6/30

60000/60000 [=====] - 5s 78us/step - loss: 0.1244 - accuracy: 0.9609 - val\_loss: 0.0800 - val\_accuracy: 0.9733

Epoch 7/30

60000/60000 [=====] - 5s 78us/step - loss: 0.1151 - accuracy: 0.9639 - val\_loss: 0.0747 - val\_accuracy: 0.9754

Epoch 8/30

60000/60000 [=====] - 4s 72us/step - loss: 0.1040 - accuracy: 0.9672 - val\_loss: 0.0735 - val\_accuracy: 0.9774

Epoch 9/30

60000/60000 [=====] - 4s 67us/step - loss: 0.1015 - accuracy: 0.9683 - val\_loss: 0.0691 - val\_accuracy: 0.9776

Epoch 10/30

60000/60000 [=====] - 4s 67us/step - loss: 0.0966 - accuracy: 0.9700 - val\_loss: 0.0658 - val\_accuracy: 0.9795

Epoch 11/30

60000/60000 [=====] - 4s 74us/step - loss: 0.0873 - accuracy: 0.9722 - val\_loss: 0.0651 - val\_accuracy: 0.9791

Epoch 12/30

60000/60000 [=====] - 4s 72us/step - loss: 0.0860 - accuracy: 0.9725 - val\_loss: 0.0653 - val\_accuracy: 0.9794

Epoch 13/30

60000/60000 [=====] - 4s 70us/step - loss: 0.0801 - accuracy: 0.9743 - val\_loss: 0.0618 - val\_accuracy: 0.9796

Epoch 14/30

60000/60000 [=====] - 4s 70us/step - loss: 0.0772 - accuracy: 0.9756 - val\_loss: 0.0616 - val\_accuracy: 0.9809

Epoch 15/30

60000/60000 [=====] - 4s 63us/step - loss: 0.0780 - accuracy: 0.9750 - val\_loss: 0.0645 - val\_accuracy: 0.9795

Epoch 16/30

60000/60000 [=====] - 4s 63us/step - loss: 0.0683 - accuracy: 0.9776 - val\_loss: 0.0615 - val\_accuracy: 0.9803

Epoch 17/30

60000/60000 [=====] - 4s 64us/step - loss: 0.0691 - accuracy: 0.9776 - val\_loss: 0.0594 - val\_accuracy: 0.9822

Epoch 18/30

60000/60000 [=====] - 4s 65us/step - loss: 0.0665 - accuracy: 0.9780 - val\_loss: 0.0563 - val\_accuracy: 0.9819

Epoch 19/30

60000/60000 [=====] - 4s 64us/step - loss: 0.0637 - accuracy: 0.9795 - val\_loss: 0.0570 - val\_accuracy: 0.9826

Epoch 20/30

60000/60000 [=====] - 4s 64us/step - loss: 0.0612 - accuracy: 0.9799 - val\_loss: 0.0563 - val\_accuracy: 0.9827

Epoch 21/30

60000/60000 [=====] - 4s 66us/step - loss: 0.0595 - accuracy: 0.9801 - val\_loss: 0.0601 - val\_accuracy: 0.9815

Epoch 22/30

60000/60000 [=====] - 4s 67us/step - loss: 0.0588 - accuracy: 0.9810 - val\_loss: 0.0556 - val\_accuracy: 0.9823

Epoch 23/30

60000/60000 [=====] - 5s 75us/step - loss: 0.0556 - accuracy: 0.9820 - val\_loss: 0.0574 - val\_accuracy: 0.9820

Epoch 24/30

60000/60000 [=====] - 5s 81us/step - loss: 0.0540 - accuracy: 0.9829 - val\_loss: 0.0523 - val\_accuracy: 0.9836

Epoch 25/30

60000/60000 [=====] - 4s 74us/step - loss: 0.0536 - accuracy: 0.9823 - val\_loss: 0.0563 - val\_accuracy: 0.9834

Epoch 26/30

```

60000/60000 [=====] - 4s 63us/step - loss: 0.0515 - accuracy: 0.9831 - va
l_loss: 0.0564 - val_accuracy: 0.9833
Epoch 27/30
60000/60000 [=====] - 4s 65us/step - loss: 0.0507 - accuracy: 0.9829 - va
l_loss: 0.0543 - val_accuracy: 0.9836
Epoch 28/30
60000/60000 [=====] - 4s 64us/step - loss: 0.0508 - accuracy: 0.9835 - va
l_loss: 0.0614 - val_accuracy: 0.9828
Epoch 29/30
60000/60000 [=====] - 4s 67us/step - loss: 0.0483 - accuracy: 0.9842 - va
l_loss: 0.0562 - val_accuracy: 0.9826
Epoch 30/30
60000/60000 [=====] - 4s 72us/step - loss: 0.0489 - accuracy: 0.9836 - va
l_loss: 0.0536 - val_accuracy: 0.9841

```

In [14]:


```

score1 = model1.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score1[0])
print('Test accuracy:', score1[1])

# list of epoch numbers
epochs = list(range(1,nb_epoch+1))
val_loss = history.history['val_loss']
train_loss = history.history['loss']
plt_epoch_vs_loss(epochs, val_loss, train_loss)

```

Test score: 0.05358278493771504  
Test accuracy: 0.9840999841690063



In [15]:

```

w_after = model1.get_weights()
print(len(w_after))

```

14

In [16]:

```

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[6].flatten().reshape(-1,1)
out_w = w_after[12].flatten().reshape(-1,1)


fig = plt.figure(figsize=(15,5))

```

```
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



## [2.2] Model2

Input(784) - ReLu(BatchNormalization(512)) - Dropout(0.5) - ReLu(BatchNormalization(256)) - Dropout(0.25) - ReLu(BatchNormalization(128)) - Dropout(0.125) - Softmax(Output(10)) - Adam Optimizer

In [17]:

```
# for relu layers we are using 'Xavier/Glorot-Normal weight Initialization'
# If we sample weights from a normal distribution  $N(0,\sigma)$  we satisfy this condition with  $\sigma=2/(\text{fan\_in}+\text{fan\_out})$ .

from keras.initializers import glorot_normal

model2 = Sequential()

model2.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=glorot_normal(
seed=None)))
model2.add(BatchNormalization())
model2.add(Dropout(0.5))

model2.add(Dense(256, activation='relu', kernel_initializer=glorot_normal(seed=None)))
model2.add(BatchNormalization())
model2.add(Dropout(0.25))

model2.add(Dense(128, activation='relu', kernel_initializer=glorot_normal(seed=None)) )
model2.add(BatchNormalization())
model2.add(Dropout(0.125))

model2.add(Dense(output_dim, activation='softmax'))

model2.summary()
```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 512)	401920



dense_4 (Dense)	(None, 512)	401328
batch_normalization_3 (Batch Normalization)	(None, 512)	2048
dropout_3 (Dropout)	(None, 512)	0
dense_5 (Dense)	(None, 256)	131328
batch_normalization_4 (Batch Normalization)	(None, 256)	1024
dropout_4 (Dropout)	(None, 256)	0
dense_6 (Dense)	(None, 128)	32896
batch_normalization_5 (Batch Normalization)	(None, 128)	512
dropout_5 (Dropout)	(None, 128)	0
dense_7 (Dense)	(None, 10)	1290
=====		
Total params: 571,018		
Trainable params: 569,226		
Non-trainable params: 1,792		

In [18]:

```
model2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model2.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/30
60000/60000 [=====] - 10s 169us/step - loss: 0.3345 - accuracy: 0.8966 - val_loss: 0.1244 - val_accuracy: 0.9605
Epoch 2/30
60000/60000 [=====] - 9s 146us/step - loss: 0.1651 - accuracy: 0.9495 - val_loss: 0.0937 - val_accuracy: 0.9716
Epoch 3/30
60000/60000 [=====] - 9s 150us/step - loss: 0.1328 - accuracy: 0.9582 - val_loss: 0.0817 - val_accuracy: 0.9732
Epoch 4/30
60000/60000 [=====] - 9s 142us/step - loss: 0.1169 - accuracy: 0.9642 - val_loss: 0.0764 - val_accuracy: 0.9767
Epoch 5/30
60000/60000 [=====] - 9s 145us/step - loss: 0.1008 - accuracy: 0.9683 - val_loss: 0.0726 - val_accuracy: 0.9774
Epoch 6/30
60000/60000 [=====] - 9s 144us/step - loss: 0.0927 - accuracy: 0.9708 - val_loss: 0.0662 - val_accuracy: 0.9799
Epoch 7/30
60000/60000 [=====] - 9s 146us/step - loss: 0.0843 - accuracy: 0.9735 - val_loss: 0.0671 - val_accuracy: 0.9800
Epoch 8/30
60000/60000 [=====] - 9s 156us/step - loss: 0.0794 - accuracy: 0.9747 - val_loss: 0.0622 - val_accuracy: 0.9807
Epoch 9/30
60000/60000 [=====] - 9s 151us/step - loss: 0.0727 - accuracy: 0.9766 - val_loss: 0.0647 - val_accuracy: 0.9787
Epoch 10/30
60000/60000 [=====] - 10s 162us/step - loss: 0.0720 - accuracy: 0.9772 - val_loss: 0.0630 - val_accuracy: 0.9805
Epoch 11/30
60000/60000 [=====] - 9s 154us/step - loss: 0.0664 - accuracy: 0.9790 - val_loss: 0.0557 - val_accuracy: 0.9833
Epoch 12/30
60000/60000 [=====] - 10s 166us/step - loss: 0.0632 - accuracy: 0.9798 - val_loss: 0.0602 - val_accuracy: 0.9823
Epoch 13/30
60000/60000 [=====] - 9s 153us/step - loss: 0.0605 - accuracy: 0.9800 - val_loss: 0.0543 - val_accuracy: 0.9846
Epoch 14/30
60000/60000 [=====] - 9s 151us/step - loss: 0.0598 - accuracy: 0.9801 - val_loss: 0.0593 - val_accuracy: 0.9825
Epoch 15/30
```

```

60000/60000 [=====] - 9s 151us/step - loss: 0.0568 - accuracy: 0.9813 - v
al_loss: 0.0586 - val_accuracy: 0.9816
Epoch 16/30
60000/60000 [=====] - 9s 155us/step - loss: 0.0512 - accuracy: 0.9835 - v
al_loss: 0.0580 - val_accuracy: 0.9840
Epoch 17/30
60000/60000 [=====] - 9s 156us/step - loss: 0.0473 - accuracy: 0.9845 - v
al_loss: 0.0540 - val_accuracy: 0.9855
Epoch 18/30
60000/60000 [=====] - 9s 156us/step - loss: 0.0481 - accuracy: 0.9847 - v
al_loss: 0.0579 - val_accuracy: 0.9826
Epoch 19/30
60000/60000 [=====] - 9s 148us/step - loss: 0.0491 - accuracy: 0.9838 - v
al_loss: 0.0476 - val_accuracy: 0.9851
Epoch 20/30
60000/60000 [=====] - 9s 145us/step - loss: 0.0440 - accuracy: 0.9856 - v
al_loss: 0.0567 - val_accuracy: 0.9845
Epoch 21/30
60000/60000 [=====] - 9s 150us/step - loss: 0.0460 - accuracy: 0.9852 - v
al_loss: 0.0595 - val_accuracy: 0.9848
Epoch 22/30
60000/60000 [=====] - 9s 152us/step - loss: 0.0415 - accuracy: 0.9862 - v
al_loss: 0.0519 - val_accuracy: 0.9848
Epoch 23/30
60000/60000 [=====] - 10s 161us/step - loss: 0.0390 - accuracy: 0.9871 - v
al_loss: 0.0527 - val_accuracy: 0.9863
Epoch 24/30
60000/60000 [=====] - 9s 146us/step - loss: 0.0406 - accuracy: 0.9872 - v
al_loss: 0.0513 - val_accuracy: 0.9850
Epoch 25/30
60000/60000 [=====] - 9s 144us/step - loss: 0.0387 - accuracy: 0.9873 - v
al_loss: 0.0493 - val_accuracy: 0.9862
Epoch 26/30
60000/60000 [=====] - 9s 146us/step - loss: 0.0375 - accuracy: 0.9875 - v
al_loss: 0.0517 - val_accuracy: 0.9857
Epoch 27/30
60000/60000 [=====] - 9s 144us/step - loss: 0.0372 - accuracy: 0.9873 - v
al_loss: 0.0502 - val_accuracy: 0.9861
Epoch 28/30
60000/60000 [=====] - 9s 144us/step - loss: 0.0359 - accuracy: 0.9884 - v
al_loss: 0.0533 - val_accuracy: 0.9853
Epoch 29/30
60000/60000 [=====] - 9s 145us/step - loss: 0.0361 - accuracy: 0.9880 - v
al_loss: 0.0500 - val_accuracy: 0.9860
Epoch 30/30
60000/60000 [=====] - 9s 145us/step - loss: 0.0341 - accuracy: 0.9885 - v
al_loss: 0.0558 - val_accuracy: 0.9841

```

In [19]:


```


score2 = model2.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score2[0])
print('Test accuracy:', score2[1])

# list of epoch numbers
epochs = list(range(1,nb_epoch+1))
val_loss = history.history['val_loss']
train_loss = history.history['loss']
plt_epoch_vs_loss(epochs, val_loss, train_loss)

```

Test score: 0.0558483196992951  
Test accuracy: 0.9840999841690063





In [20]:

```
w_after = model2.get_weights()
print(len(w_after))
```

20

In [21]:

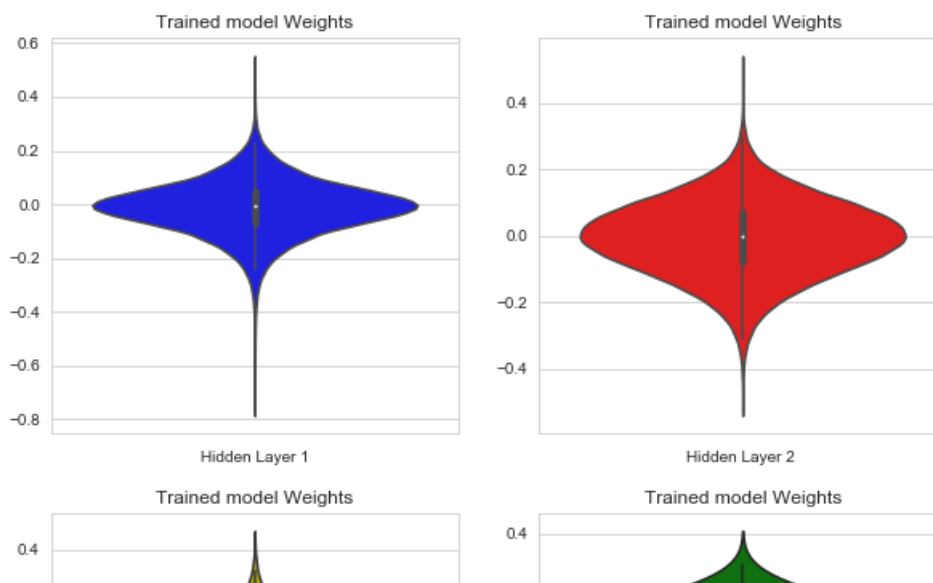
```
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[6].flatten().reshape(-1,1)
h3_w = w_after[12].flatten().reshape(-1,1)
out_w = w_after[18].flatten().reshape(-1,1)


fig = plt.figure(figsize=(10,10))
plt.title("Weight matrices after model trained")
plt.subplot(2, 2, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(2, 2, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')


plt.subplot(2, 2, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w,color='y')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(2, 2, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='g')
plt.xlabel('Output Layer ')
plt.show()
```





Hidden Layer 3



Output Layer

## [2.3] Model3

Input(784) - ReLu(BatchNormalization(512)) - Dropout(0.5) - ReLu(BatchNormalization(256)) - Dropout(0.4) - ReLu(BatchNormalization(128)) - Dropout(0.3) - ReLu(BatchNormalization(64)) - Dropout(0.2) - ReLu(BatchNormalization(32)) - Dropout(0.1) - Softmax(Output(10)) - Adam Optimizer

In [22]:

```
# for relu layers we are using 'He-Normal weight Initialization'
# If we sample weights from a normal distribution  $N(0,\sigma)$  we satisfy this condition with  $\sigma=2/(\text{fan\_in}+\text{fan\_out})$ .

from keras.initializers import he_normal

model3 = Sequential()

model3.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal(seed=None)))
model3.add(BatchNormalization())
model3.add(Dropout(0.5))

model3.add(Dense(256, activation='relu', kernel_initializer=he_normal(seed=None)))
model3.add(BatchNormalization())
model3.add(Dropout(0.4))

model3.add(Dense(128, activation='relu', kernel_initializer=he_normal(seed=None)))
model3.add(BatchNormalization())
model3.add(Dropout(0.3))

model3.add(Dense(64, activation='relu', kernel_initializer=he_normal(seed=None)))
model3.add(BatchNormalization())
model3.add(Dropout(0.2))

model3.add(Dense(32, activation='relu', kernel_initializer=he_normal(seed=None)))
model3.add(BatchNormalization())
model3.add(Dropout(0.1))

model3.add(Dense(output_dim, activation='softmax'))

model3.summary()
```

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
dense_8 (Dense)	(None, 512)	401920
batch_normalization_6 (Batch Normalization)	(None, 512)	2048
dropout_6 (Dropout)	(None, 512)	0
dense_9 (Dense)	(None, 256)	131328
batch_normalization_7 (Batch Normalization)	(None, 256)	1024
dropout_7 (Dropout)	(None, 256)	0
dense_10 (Dense)	(None, 128)	32896

batch_normalization_8 (Batch Normalization)	(None, 128)	512
dropout_8 (Dropout)	(None, 128)	0
dense_11 (Dense)	(None, 64)	8256
batch_normalization_9 (Batch Normalization)	(None, 64)	256
dropout_9 (Dropout)	(None, 64)	0
dense_12 (Dense)	(None, 32)	2080
batch_normalization_10 (Batch Normalization)	(None, 32)	128
dropout_10 (Dropout)	(None, 32)	0
dense_13 (Dense)	(None, 10)	330
=====		
Total params: 580,778		
Trainable params: 578,794		
Non-trainable params: 1,984		

In [23]:

```
model3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model3.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/30
60000/60000 [=====] - 12s 194us/step - loss: 0.6755 - accuracy: 0.7921 - val_loss: 0.1842 - val_accuracy: 0.9442
Epoch 2/30
60000/60000 [=====] - 10s 168us/step - loss: 0.2663 - accuracy: 0.9239 - val_loss: 0.1292 - val_accuracy: 0.9639
Epoch 3/30
60000/60000 [=====] - 10s 166us/step - loss: 0.2076 - accuracy: 0.9409 - val_loss: 0.1088 - val_accuracy: 0.9686
Epoch 4/30
60000/60000 [=====] - 10s 167us/step - loss: 0.1793 - accuracy: 0.9498 - val_loss: 0.0968 - val_accuracy: 0.9723
Epoch 5/30
60000/60000 [=====] - 11s 175us/step - loss: 0.1575 - accuracy: 0.9552 - val_loss: 0.0918 - val_accuracy: 0.9719
Epoch 6/30
60000/60000 [=====] - 10s 167us/step - loss: 0.1455 - accuracy: 0.9595 - val_loss: 0.0854 - val_accuracy: 0.9745
Epoch 7/30
60000/60000 [=====] - 11s 182us/step - loss: 0.1315 - accuracy: 0.9624 - val_loss: 0.0867 - val_accuracy: 0.9758
Epoch 8/30
60000/60000 [=====] - 12s 203us/step - loss: 0.1258 - accuracy: 0.9642 - val_loss: 0.0778 - val_accuracy: 0.9780
Epoch 9/30
60000/60000 [=====] - 10s 162us/step - loss: 0.1154 - accuracy: 0.9670 - val_loss: 0.0777 - val_accuracy: 0.9770
Epoch 10/30
60000/60000 [=====] - 10s 165us/step - loss: 0.1065 - accuracy: 0.9694 - val_loss: 0.0750 - val_accuracy: 0.9790
Epoch 11/30
60000/60000 [=====] - 11s 191us/step - loss: 0.1050 - accuracy: 0.9700 - val_loss: 0.0702 - val_accuracy: 0.9806
Epoch 12/30
60000/60000 [=====] - 12s 201us/step - loss: 0.0978 - accuracy: 0.9717 - val_loss: 0.0694 - val_accuracy: 0.9804
Epoch 13/30
60000/60000 [=====] - 12s 195us/step - loss: 0.0936 - accuracy: 0.9726 - val_loss: 0.0739 - val_accuracy: 0.9802
Epoch 14/30
60000/60000 [=====] - 11s 191us/step - loss: 0.0907 - accuracy: 0.9739 - val_loss: 0.0628 - val_accuracy: 0.9823
Epoch 15/30
60000/60000 [=====] - 11s 184us/step - loss: 0.0850 - accuracy: 0.9754 -
```

```

00000/00000 [-----] - 11s 104us/step - loss: 0.0800 - accuracy: 0.9754 -
val_loss: 0.0624 - val_accuracy: 0.9807
Epoch 16/30
60000/60000 [=====] - 11s 187us/step - loss: 0.0804 - accuracy: 0.9764 -
val_loss: 0.0683 - val_accuracy: 0.9807
Epoch 17/30
60000/60000 [=====] - 11s 185us/step - loss: 0.0807 - accuracy: 0.9772 -
val_loss: 0.0645 - val_accuracy: 0.9812
Epoch 18/30
60000/60000 [=====] - 11s 176us/step - loss: 0.0772 - accuracy: 0.9767 -
val_loss: 0.0604 - val_accuracy: 0.9834
Epoch 19/30
60000/60000 [=====] - 10s 174us/step - loss: 0.0726 - accuracy: 0.9789 -
val_loss: 0.0597 - val_accuracy: 0.9830
Epoch 20/30
60000/60000 [=====] - 11s 177us/step - loss: 0.0701 - accuracy: 0.9797 -
val_loss: 0.0562 - val_accuracy: 0.9842
Epoch 21/30
60000/60000 [=====] - 10s 167us/step - loss: 0.0683 - accuracy: 0.9807 -
val_loss: 0.0602 - val_accuracy: 0.9833
Epoch 22/30
60000/60000 [=====] - 10s 164us/step - loss: 0.0695 - accuracy: 0.9800 -
val_loss: 0.0603 - val_accuracy: 0.9831
Epoch 23/30
60000/60000 [=====] - 10s 174us/step - loss: 0.0659 - accuracy: 0.9807 -
val_loss: 0.0636 - val_accuracy: 0.9825
Epoch 24/30
60000/60000 [=====] - 11s 184us/step - loss: 0.0648 - accuracy: 0.9810 -
val_loss: 0.0634 - val_accuracy: 0.9832
Epoch 25/30
60000/60000 [=====] - 12s 202us/step - loss: 0.0639 - accuracy: 0.9810 -
val_loss: 0.0599 - val_accuracy: 0.9836
Epoch 26/30
60000/60000 [=====] - 12s 202us/step - loss: 0.0623 - accuracy: 0.9818 -
val_loss: 0.0601 - val_accuracy: 0.9827
Epoch 27/30
60000/60000 [=====] - 11s 190us/step - loss: 0.0614 - accuracy: 0.9823 -
val_loss: 0.0577 - val_accuracy: 0.9843
Epoch 28/30
60000/60000 [=====] - 12s 197us/step - loss: 0.0568 - accuracy: 0.9837 -
val_loss: 0.0547 - val_accuracy: 0.9852
Epoch 29/30
60000/60000 [=====] - 12s 196us/step - loss: 0.0566 - accuracy: 0.9831 -
val_loss: 0.0600 - val_accuracy: 0.9846
Epoch 30/30
60000/60000 [=====] - 12s 206us/step - loss: 0.0541 - accuracy: 0.9840 -
val_loss: 0.0531 - val_accuracy: 0.9858

```

In [24]:

```


score3 = model3.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score3[0])
print('Test accuracy:', score3[1])


# list of epoch numbers
epoch = list(range(1,nb_epoch+1))
val_loss = history.history['val_loss']
train_loss = history.history['loss']
plt_epoch_vs_loss(epoch, val_loss, train_loss)

```

Test score: 0.053111556177656165

Test accuracy: 0.98580002784729





In [25]:

```
w_after = model3.get_weights()
print(len(w_after))
```

32

In [26]:

```
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[6].flatten().reshape(-1,1)
h3_w = w_after[12].flatten().reshape(-1,1)
h4_w = w_after[18].flatten().reshape(-1,1)
h5_w = w_after[24].flatten().reshape(-1,1)
out_w = w_after[30].flatten().reshape(-1,1)

fig = plt.figure(figsize=(15,10))
plt.title("Weight matrices after model trained")
plt.subplot(2, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')




plt.subplot(2, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

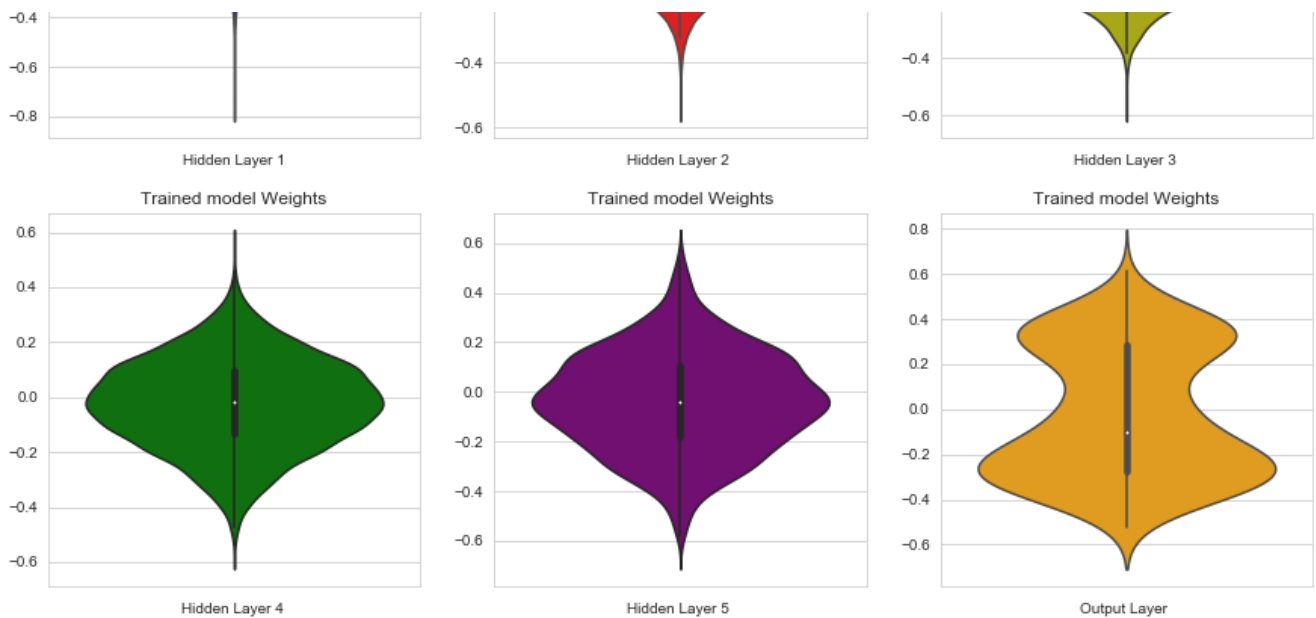
plt.subplot(2, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w,color='y')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(2, 3, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w,color='g')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(2, 3, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w,color='purple')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(2, 3, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='orange')
plt.xlabel('Output Layer ')
plt.show()
```





### [3] Results

In [27]:

```
from prettytable import PrettyTable

table = PrettyTable()
table.field_names = ["Model", "Hidden Layers", "Score", "Accuracy"]
table.add_row([1, 2, round(score1[0], 3), round(score1[1], 3)])
table.add_row([2, 3, round(score2[0], 3), round(score2[1], 3)])
table.add_row([3, 5, round(score3[0], 3), round(score3[1], 3)])

print(table.get_string(title="Results"))
```

Results				
Model	Hidden Layers	Score	Accuracy	
1	2	0.054	0.984	
2	3	0.056	0.984	
3	5	0.053	0.986	

### [4] Conclusion

There is no much difference in the accuracy of all models.